

Report of training the agent

The model was trained for a maximum of 2,000 episodes but the agent was able to solve the environment (i.e. get atleast +13 average score over 100 adjacent episodes).

Model

The model is called a Deep Q-Network which was first discovered by the DeepMind team. This uses a Deep Neural Network (or a Multi Layered Perceptron) for estimating the action-value function (Q_{π}).

The Q-Network has three dense (or fully connected layers). The first two layers have 64 nodes activated with ReLU activation function. The final (output layer) has 4 nodes and is activated with linear activation (or no activation at all). This network takes in as input the 37 dimensional current state and gives as output 4 action-values corresponding to the possible actions that the agent can take.

The neural network used Adam optimizer and Mean Squared Error (MSE) as the loss function.

The following image provides a pictorial representation of the Q-Network model:



The following image provides the plot for score v/s episode number:



Hyperparameters used

Hyperparameter	Value	Description
Buffer size	100000	Maximum size of the replay buffer
Batch size	64	Batch size for sampling from replay buffer
Gamma (γ)	0.99	Discount factor for calculating return
Tau (τ)	0.001	Hyperparameter for soft update of target parameters
Learning Rate (α)	0.0005	Learning rate for the neural networks
Update Every (C)	4	Number of time steps after which soft update is performed
Epsilon (ϵ)	1.0	For epsilon-greedy action selection
Epsilon decay rate	0.995	Rate by which epsilon decays after every episode
Epsilon minimum	0.01	The minimum value of epsilon

Future work

The following algorithms can be considered for further development of this agent:

- Double DQN
- Duelling DQN
- Prioritized Experience Replay DQN
- Rainbow